


GPU 显存支持按照 MiB 分配资源

1. 背景

使用阿里云的插件，<https://github.com/AliyunContainerService/gpushare-scheduler-extender>

安装插件的 yaml 中，显存的使用是 GiB 为单位

```
unipolicy: ClusterPIPSwitchHostNet
containers:
  #- image: registry.cn-hangzhou.aliyuncs.com/acs/k8s-gpushare-p
  - image: docker.dm-ai.cn/public/gpushare-device-plugin:v0.1.3
    name: gpushare
    # Make this pod as Guaranteed pod which will never be evicted
    command:
      - gpushare-device-plugin-v2
      - -logtostderr
      - --v=5
      - --memory-unit=GiB
    resources:
      limits:
        memory: "300Mi"
        cpu: "1"
      requests:
        memory: "300Mi"
        cpu: "1"
    env:
      - name: KUBECONFIG
```



此时资源申请如下，表示使用 4GB 显存

```

- name: APOLLO_CLUSTER_NAME
  value: default
- name: APOLLO_NAMESPACE
  value: application
resources:
  limits:
    aliyun.com/gpu-mem: '4'
    cpu: '7'
    memory: 6000Mi
  requests:
    cpu: '5'
    memory: 5000Mi
  livenessProbe:
    tcpSocket:
      port: 80
    initialDelaySeconds: 15
    timeoutSeconds: 1
    periodSeconds: 20

```

查看显存大小时，也是以 GiB 为单位，效果类似如下：

```

[root@jenkins ~]# kubectl inspect gpushare
NAME      IPADDRESS  GPU0(Allocated/Total) GPU1(Allocated/Total) GPU2(Allocated/Total) GPU3(Allocated/Total) GPU4(Allocated/Total) GPU5(Allocated/Total) GPU6(Allocated/Total) GPU7(Allocated/Total) GPU Memory (GiB)
192.168.68.13 192.168.68.13 8/11 8/11 8/11 8/11 8/11 8/11 8/11 8/11 64/88
Allocated/Total GPU Memory In Cluster:
64/88 (220)

```

2. 变更

将脚本中 GiB 改为 MiB 后

```

command:
- gpushare-device-plugin-v2
- -logtostderr
- --v=5
- --memory-unit=MiB
resources:
  limits:
    memory: "300Mi"
    cpu: "1"
  requests:
    memory: "300Mi"
    cpu: "1"

```

无法识别显卡，对应节点的 kubelet 报错：

```

pectedly for device plugin aliyun.com/gpu-mem with error rpc error: code = ResourceExhausted desc = gr
pc: received message larger than max (5365984 vs. 4194304)
Apr  5 15:00:27 node28 kubelet: E0405 15:00:27.586718 35954 endpoint.go:106] listAndwatch ended unex
pectedly for device plugin aliyun.com/gpu-mem with error rpc error: code = ResourceExhausted desc = gr
pc: received message larger than max (5365984 vs. 4194304)
Apr  5 15:03:18 node28 kubelet: E0405 15:03:18.297762 37987 endpoint.go:106] listAndwatch ended unex
pectedly for device plugin aliyun.com/gpu-mem with error rpc error: code = ResourceExhausted desc = gr
pc: received message larger than max (5365984 vs. 4194304)
Apr  5 15:23:58 node28 kubelet: E0405 15:23:58.192347 9900 endpoint.go:106] listAndwatch ended unex
pectedly for device plugin aliyun.com/gpu-mem with error rpc error: code = ResourceExhausted desc = gr
pc: received message larger than max (5365984 vs. 4194304)
[root@k8s-192-168-68-5 ~]#

```

Apr 5 15:03:18 node28 kubelet: E0405 15:03:18.297762 37987 endpoint.go:106] listAndwatch ended unexpectedly for device plugin aliyun.com/gpu-mem with error rpc error: code = ResourceExhausted desc = grpc: received message larger than max (5365984 vs. 4194304)

这是因为 GRPC 源码中，接收的 grpc 消息最大报文是 $1024 \times 1024 \times 4 = 4\text{MiB}$ ，消息体超过了最大大小 4MB。

3. 解决办法

重新编译 kubelet 源码，替换原有 kubelet。修改代码的地方是

<https://github.com/kubernetes/kubernetes/blob/master/pkg/kubelet/cm/devicemanager/endpoint.go> 中，

dial 方法中添加参数

```
grpc.WithDefaultCallOptions(grpc.MaxCallRecvMsgSize(1024*1024*16))
```

如下图：

```

func dial(unixSocketPath string) (pluginapi.DevicePluginClient, *grpc.ClientConn, error) {
    ctx, cancel := context.WithTimeout(context.Background(), 10*time.Second)
    defer cancel()

    c, err := grpc.DialContext(ctx, unixSocketPath, grpc.WithInsecure(), grpc.WithBlock(),
        grpc.WithContextDialer(func(ctx context.Context, addr string) (net.Conn, error) {
            return (&net.Dialer{}).DialContext(ctx, "unix", addr)
        }), grpc.WithDefaultCallOptions(grpc.MaxCallRecvMsgSize(1024*1024*16)),
    )

    if err != nil {
        return nil, nil, fmt.Errorf(errFailedToDialDevicePlugin+" %v", err)
    }

    return pluginapi.NewDevicePluginClient(c), c, nil
}

```

pkg/kubelet/cm/devicemanager/endpoint.go" 190L, 5562C

重新编译 kubelet

make WHAT=cmd/kubelet

```
[root@devops-10-12-19-31 kubernetes-1.18.10]# make WHAT=cmd/kubelet
+++ [0521 16:57:53] Building go targets for linux/amd64:
./vendor/k8s.io/code-generator/cmd/deepcopy-gen
+++ [0521 16:58:02] Building go targets for linux/amd64:
./vendor/k8s.io/code-generator/cmd/defaulter-gen
+++ [0521 16:58:10] Building go targets for linux/amd64:
./vendor/k8s.io/code-generator/cmd/conversion-gen
+++ [0521 16:58:24] Building go targets for linux/amd64:
./vendor/k8s.io/kube-openapi/cmd/openapi-gen
+++ [0521 16:58:36] Building go targets for linux/amd64:
./vendor/github.com/go-bindata/go-bindata/go-bindata
warning: ignoring symlink /data/kubernetes-1.18.10/_output/local/go/src/k8s.io/kubernetes
go: warning: "k8s.io/kubernetes/vendor/github.com/go-bindata/go-bindata/..." matched no packages
+++ [0521 16:58:38] Building go targets for linux/amd64:
cmd/kubelet
[root@devops-10-12-19-31 kubernetes-1.18.10]#
```

编译成功的二进制文件路径：_output/bin/kubelet

将生成的 kubelet 二进制文件进行替换，重启 kubelet

4. 效果

上述步骤搞完后，查看 gpu 情况：

```
[root@jenkins ~]# kubectl inspect gpushare
NAME      IPADDRESS  GPU0(allocated/Total) GPU1(allocated/Total) GPU2(allocated/Total) GPU3(allocated/Total) GPU4(allocated/Total) GPU5(allocated/Total) GPU6(allocated/Total) GPU7(allocated/Total) GPU Memory(
192.168.68.13 192.168.68.13 8/12066 8/12066 8/12066 8/12066 8/12066 8/12066 8/12066 8/12066 64/96528
Alllocated/Total GPU Memory in cluster:
64/96528 (0%)
```